

Ashish Sahoo

Hyderabad, India | ashish.sahoo12@gmail.com | +91-7008897871 | <https://www.linkedin.com/in/ashishsahoo07/>

SUMMARY

Senior Data Engineer with 7 years of experience in SQL, Python, and ETL development, delivering scalable data solutions and automating data workflows across on-premise and cloud platforms. Expertise in risk management and regulatory reporting, with experience building complex data models for financial data processing. Proven ability to ensure data accuracy, governance, and compliance across large-scale pipelines. Currently leading a team delivering reliable, high-performance, audit-ready data solutions.

TECHNICAL SKILLS

Programming Languages: Python, SQL, Java

Databases: MS SQL, MySQL, Oracle Database

Big Data and Processing: PySpark, Spark SQL, Delta Lake, Hive, Data Warehousing

Cloud: AWS (Lambda, S3, Glue, RDS, AppSync), Azure (ADF, Databricks, Data Lake, Azure DevOps), Snowflake

Data Engineering: Data Lake and Medallion Architecture, ETL Pipelines, Data Modeling, Data Governance

Tools and Scheduler: ODI, SSIS, Informatica Cloud, Tableau, Power BI, Azure DevOps (CI/CD), GitHub, Jira

Certification: Microsoft Azure Fundamentals (AZ-900) certified

EXPERIENCE

Senior Data Engineer

July 2024 - Present

Goldman Sachs | Hyderabad, India

- Led a team to design and deliver scalable ETL pipelines, data warehousing solutions, and complex data models using Python, SQL, PySpark and Snowflake, supporting enterprise-scale regulatory reporting systems and improving data accessibility and performance.
- Performed SQL optimization and database performance analysis on DB2 (application database), and automated reporting workflows using Python and ETL processes, reducing manual effort by ~80% and improving efficiency.
- Developed Tableau dashboards and collaborated with cross-functional stakeholders including audit, risk, and reporting teams to deliver solutions across the end-to-end SDLC using Agile methodologies.
- Developed Python-based automation using pandas, requests, and REST APIs to orchestrate LLM API calls, parse structured outputs, validate JSON payloads, and load processed data into enterprise data platforms, improving pipeline efficiency and supporting AI-driven workflows.
- Led migration of legacy data models to the Alloy platform, developing reusable REST-based Alloy Services and enabling row-level entitlement models for secure semantic-layer access.

Senior Data Engineer (Specialist Programmer)

Jan 2023 - May 2024

Infosys Limited | Czech Republic, Europe

- Worked onsite at client location in the Czech Republic (Europe), providing production support and monitoring pipelines for high availability, resolving SSIS-based data integration issues via JIRA, and leading migration of legacy on-prem SSIS solutions to Azure cloud services.
- Developed and tested ADF and Azure Databricks pipelines integrating data from Oracle Database, Mainframe, and SAP, using Spark SQL transformations to ensure data quality and consistency.
- Automated deployments using 10+ Azure DevOps YAML pipelines, optimizing CI/CD for cleaned and curated data layers.

Data Engineer (Specialist Programmer)

July 2019 - Dec 2022

Infosys Limited | Hyderabad, India

- Built scalable ETL pipelines using Oracle Data Integrator (ODI) and Informatica IICS, reducing operational costs by ~30%, and automated file archival, profiling, and ingestion workflows using Python.
- Engineered AWS-based data pipelines using Glue (PySpark), Lambda (Node.js), Aurora PostgreSQL, and AWS AppSync, enabling seamless data migration and API-driven consumption.
- Transformed and cleansed historical data using PostgreSQL (AWS RDS), ensuring structured and reliable datasets for downstream systems.

EDUCATION

Bachelor of Engineering in Computer Science

July 2015 - May 2019

Institute of Technical Education and Research, SOA University, India